

Supplementary material: Graph algorithms for predicting subcellular localization at the pathway level

Chris S Magnano and Anthony Gitter

Supplementary Methods

5.1. *Experimental Setup*

5.1.1. *Spatial Proteomics Case Study*

Mass Spectrometry Datasets: To investigate a possible use-case of pathway-based localization prediction, and validate its performance, we performed localization prediction on proteomic quantification of primary fibroblasts during human cytomegalovirus (HCMV) infection. Two mass spectrometry quantification methods were used at 5 time points post infection: 24, 48, 72, 96, and 120 hours post infection (hpi). The first of these datasets provides label-free protein quantification at each timepoint. The other was quantified using isobaric labeling via tandem mass tags (TMTs). These two datasets will be referred to as the label-free and the TMT datasets, respectively.

Multi-organelle profiling was performed on the TMT dataset via gradient centrifugation to fractionate organelles. This process partially separates organelles into a set of subcellular fractions. While each of these fractions are not purely a single organelle, each organelle contains a unique signature in its quantification across the subcellular fractions. Clustering analysis is then used to group proteins with similar fraction profiles. Each cluster is labeled as a certain organelle via a set of marker proteins; proteins known with high-confidence to localize to a particular organelle.

In the HCMV protein quantification a set of marker proteins were curated from UniProt subcellular location annotations with experimental evidence; proteins that were annotated with multiple localizations were excluded from the marker set. Proteins that were not confidently assigned to a particular organelle were left as unlabeled. There were 2730 proteins in total, of which 1229 had localization labels at 24hpi and 1348 had labels at 120hpi. 574 of these proteins were marker proteins.

Experimental Scenarios: We explored 3 different data availability scenarios for using the best performing models from the pathway database prediction task, GAT, in an experimental setting. For all 3 scenarios, the best performing hyperparameter combination from the pathway database prediction task was used, no further hyperparameter tuning was performed. The 3 scenarios are as follows:

- (1) First, we trained a model using data from the PathBank database as described in Section 2.1. Here we took the best performing pre-trained model with no further modifications and used it to predict localizations at the 120hpi timepoint.
- (2) Second, we trained a model using a separate dataset which measured protein localization using a similar method on a different cell type and under a different biological condition,

here HeLa cells undergoing EGF stimulation.

- (3) Third, we trained a model on the same experiment at the 24hpi timepoint. This third scenario is unlikely to occur, as it would require a dataset to already exist for an identical cell type and condition, but gives a useful benchmark for predictive performance.

Pathway reconstruction: Pathway reconstruction was performed using OmicsIntegrator2 ^d, which is similar to Omics Integrator but uses a different optimization algorithm. OmicsIntegrator2 performs pathway reconstruction via the prize-collecting Steiner forest problem. Pathway reconstruction was run using the Signaling Pathway Reconstruction Analysis Streamliner (SPRAS) ^e. Protein fold-change was used as prizes for both HCMV fibroblast infection data and HeLa EGF stimulation data. The set of all protein protein interactions used as the background set of candidate edges for pathway reconstruction consists of a previously used network with 161,901 weighted edges, based on merged protein interactions from the iRefIndex database v13 ^f and kinase-substrate interactions from PhosphoSitePlus ^g. ω was tested between 1 and 10, β was tested between 1 and 5, and μ was tested between 0.1 and 1. Each hyperparameter was evaluated at 10 increments across its range. For the 120hpi timepoint pathway parameter advising was used to select the top 50 pathways from 1000 candidate hyperparameter combinations. For the 24hpi timepoint and HeLa stimulation dataset, all nonempty pathways were used so that these pathway datasets were comparable in size to the pathway database datasets. All models hyperparameters were the same as those used in the highest performing GAT model, using Compartments features to predict PathBank labels, as recorded in Table S2.

5.2. Models

5.2.1. Neural Networks

The maximum value in the output layer is used as the final classification. All neural network models were trained using cross-entropy loss. All neural networks were implemented using the PyTorch Geometric package.

Graph convolutional network: The graph convolutional network incorporated a set of message-passing convolutional layers before the final set of fully connected layers. These convolutional layers allow for information to be shared across the topology of the input network. The l^{th} convolutional layer $H^{(l)}$ are updated via the following rule:

$$H^{(l)} = ReLU(D^{-\frac{1}{2}} \tilde{A} D^{\frac{1}{2}} H^{(l-1)} W^{(l-1)})$$

Where \tilde{A} is the adjacency matrix of the input pathway with added self-edges for all nodes, D is a degree matrix normalization factor where $D_{ii} = \sum_j \tilde{A}_{ij}$, and $W^{(l)}$ is a set of weights for the

^d<https://github.com/fraenkel-lab/OmicsIntegrator2>

^e<https://github.com/Reed-CompBio/spras>

^f<https://irefindex.vib.be/wiki/index.php/iRefIndex>

^g<https://www.phosphosite.org>

l^{th} layer. This update rule provides a first-order approximation of spectral graph convolutions and is implemented in the *GCNConv* class in PyTorch Geometric.

Graph attention network: Graph attention networks extend graph convolutional networks by allowing each node choose which of its neighbors to pay attention to. As opposed to taking the average of its neighbors, each node computes a weighted average of its neighbors in graph convolutional layers. Furthermore, many attention networks are multi-headed, where multiple attention weights are computed in parallel for each node. The number of heads to include is an input hyperparameter, and generally increases accuracy at the cost of increased computational complexity. We used the *GATV2Conv* class for graph layers, which is a more expressive implementation of graph attention networks that allows for more diversity in attention between nodes.

Graph isomorphism network: Graph isomorphism networks take advantage of the similarity between neighbor aggregation in graph neural networks and the Weisfeiler-Lehman (WL) graph isomorphism test. The WL graph isomorphism test is a heuristic algorithm for determining graph isomorphisms. For two graphs, each iteration of the test every node aggregates its neighbors into a unique hash. These hashes are compared between the two graphs, and if they differ the graphs are known to be non-isomorphic. Iterations of the test are repeated until the user feels confident that the graphs are isomorphic; the algorithm cannot conclusively prove isomorphism.

The neighbor aggregation in each graph layer of a graph isomorphism network is formulated to be at least as powerful as the WL isomorphism test; the l^{th} layer is guaranteed to generate different embeddings of two graphs if those graphs would be found to be non-isomorphic via the WL isomorphism test in l iterations. The representation of each node in layer l of a graph isomorphism network, $h_n^{(l)}$, is computed as:

$$h_n^{(l)} = MLP^{(l)}((1 + \epsilon^{(l)}) \cdot h_n^{(l-1)} + \sum_{u \in Adj(n)} h_u^{(l-1)})$$

Where *MLP* is a multi-layer perceptron, ϵ is a learned parameter, and *Adj*(n) is the set of nodes adjacent to n in the input pathway. We used the *GINConv* class in PyTorch Geometric for graph isomorphism layers.

5.2.2. Probabilistic Graphical Models

Given the nature of the label propagation inherent in a pathway level localization prediction algorithm, and that many localization databases contain uncertain or even probabilistic data, probabilistic graphical models are a natural choice for an initial approach to pathway level localization prediction. As moving between subcellular locations costs energy, it is unlikely to happen often within a single pathway. Therefore, we can make the assumption from a modeling perspective that the subcellular location of an interaction is dependent on the subcellular location of neighboring interactions within a pathway.

Probabilistic graphical models represent a set of N random variables \mathbf{y} as nodes and de-

dependencies between them as a set of edges E . We created two pairwise undirected probabilistic graphical models, which we call NaivePGM and TrainedPGM. In these probabilistic graphical models the random variables obey a local Markov property, such that each random variable is conditionally independent of all others given its neighbors in the graph.

The NaivePGM is a Markov random field, where the joint probability of all localizations can be factorized as

$$P(\mathbf{y}) = \frac{1}{Z} \prod_{i \in N} \phi_i(y_i) \prod_{i,j \in E} \phi_{ij}(y_i, y_j)$$

Where Z is a normalizing function so that the combination of all possible configurations of \mathbf{y} sum to 1 and $\phi_i y_i$ and $\phi_{ij} y_i, y_j$ are the unary and pairwise potential functions, respectively. The unary potential function defines the probabilities of a given node having each localization, while the pairwise potential functions define the joint probability of each pair of nodes that share an edge. For finding the task of finding the most likely configuration of \mathbf{y} , referred to as decoding, Z can be ignored. In the NaivePGM the input features are not used to parameterize the potential functions. Instead, the unary potential functions directly map the normalized features to class probabilities, and the joint probability tables directly map the normalized features to join probability tables. This was chosen as both the ComPPI and Compartments scores represent confidences, with ComPPI scores directly representing probabilities for each localization.

The TrainedPGM is a conditional random field where the input features are treated as observations of additional variables. The probability of localization assignments \mathbf{y} are then conditioned over the input features \mathbf{x} as:

$$P(\mathbf{y}|\mathbf{x}) = \frac{1}{Z} \prod_{i \in N} \phi(y_i, x_i) \prod_{i,j \in E} \phi(x_i, x_j)$$

Here, the unary potential functions are now conditioned on observations of features x_i corresponding to each random variable y_i . Pairwise potentials are a function of observations, thus, each random variable y_i is conditionally independent of all other variables given its corresponding features/observations x_i .

The potential functions in conditional random fields are typically log-linear functions of the form $e^{\mathbf{w}_i^T \phi_f(x_i, y_i)}$, parameterized via a weight vector \mathbf{w} and ϕ_f simply represents features for each node. Additionally, typically the feature weight vectors are shared between nodes or sets of nodes. Thus, the entire model can then be represented as:

$$p(\mathbf{y}|\mathbf{x}, \mathbf{w}) = \frac{1}{Z} \exp \left(\sum_{i \in N} \mathbf{w}_f^T \phi_f(x_i, y_i) + \sum_{i,j \in E} \mathbf{w}_e^T \phi_e(x_i, x_j) \right)$$

Where $\phi_f(x_i, y_i)$ is the single unary potential function that represents features for each node, and $\phi_e(x_i, x_j)$ is the single pairwise potential function that represents combinations of features. The weight vector \mathbf{w}_f is a set of weights for each feature to each possible configuration

of y_i , while the weight vector \mathbf{w}_e is a set of weights for each feature and combination of configurations for y_i and y_j .

When represented with these potentials, the log likelihood of the model parameters \mathbf{w} can be easily represented, and is differentiable, allowing for parameters to be learned by maximum likelihood estimation via gradient-based optimization. Sets of nodes and edges can share the same set of model parameters, referred to as parameter tying.

However, parameter learning for a conditional random field of this form did not converge when trained with stochastic gradient descent. This may be due to the underlying label distributions of different pathways being too different from each other.

Instead a model was chosen where potentials are represented by discriminative classifiers, here random forests. This type of model is referred to as a discriminative random field. This was chosen over a more traditional log linear parameterization due to better performance on validation data during model selection.

These models only provide predictions on the nodes of the graph, while we are interested in localization labels on the edges. To convert the input pathway into an appropriate graphical model, each pathway is converted into a bipartite graph, where a node is added that represents each edge. First, all nodes from the original pathway are added to the graphical model. No edges from the original pathway are added. Instead, for each edge e_{ij} between nodes n_i and n_j , a node $n_{e_{ij}}$ is added representing the interaction. Then 2 edges are added to the graphical model going from each original node to the new interaction node, $n_{ie_{ij}}$ and $n_{je_{ij}}$. An overview of this process can be found in FigureS1. Each of the K features f_{ek} in $n_{e_{ij}}$ are computed as the normalized product of features from n_i and n_j , here represented as f_{ik} and f_{jk} :

$$f_{ek} = \frac{f_{ik}f_{jk}}{\sum_{l=1}^K f_{el}}$$

This is equivalent to how interaction localization probabilities are calculated in ComPPI. Parameters are tied such that all original nodes are represented by one set of model parameters, and all interaction nodes are represented by another. This can be seen in panel C of Figure S1, where ϕ_1 is the set of model parameters that describes the relationship between input features for each protein and its localization, ϕ_2 describes the relationship between each interaction's combined features and its localization, and ϕ_3 describes the relationship between each protein and the interactions it participates in.

Final localization labels can be viewed as a maximum a posteriori (MAP) estimate of the configuration of all interaction node labels. Decoding was performed using loopy belief propagation, which approximates the MAP estimate via a message passing algorithm. Loopy belief propagation was run for 10000 iterations in all cases. Both models were implemented in the DGM software library ^h.

^h<https://research.project-10.de/dgmdoc/index.html>

5.2.3. Other Classification Models

Two non-neural network classifiers were used to further examine the effect of incorporating topological information into localization prediction: logistic regression, referred to as Logit, and random forests, referred to as RF. Interactions were represented by concatenating the features of the 2 nodes that make up that interaction, in alphanumeric order by protein identifier as shown in Figure S2. Tested hyperparameter ranges for these models are listed in Table S1. Both models were implemented in Python 3.9 using the Scikit-Learn package v1.0.2.

5.3. Data

5.3.1. Pathway Databases

Pathway datasets were constructed from the Reactome and PathBank databases. Pathways were downloaded from Pathway Commons, and localization information was retrieved from BioPax pathway representations using the PyBioPax package v0.1.0ⁱ. Reactome contains localization information for all edges. PathBank, however, contains nodes and edges with missing localization information. These missing data were excluded from all analyses.

The original pathways in both Reactome and PathBank are represented as hypergraphs, where reaction edges can contain more than two nodes. PathwayCommons converts these hypergraphs to graphs using different rules for different types of reactions^j. This hypergraph conversion can affect the resulting graph topology and class distribution. During preliminary experiments we found that the hypergraph conversion resulted in an over-representation of protein-complex edges. To represent a protein-complex that contains n proteins, the hypergraph conversions creates an edge between every possible pair of nodes, resulting in n^2 edges. For instance, in PathBank the pathway Protein Synthesis: Serine as a hypergraph has 4 hyperedges. However, when converted to a graph the pathway contains 3318 edges, of which 3315 are of type “in-complex-with”. In Reactome after hypergraph conversion over 75% of edges among all pathways are of type “in-complex-with”. While this preserves protein information, for an edge classification task this conversion results in highly skewed data.

We collapsed protein complexes into single nodes where possible in all pathways. This was done by removing any node from a pathway if all of its edges were redundant with the protein-complex’s edges. A single node for each complex was left in the pathway to represent that complex. This significantly reduced the size of some pathways. For instance, the Protein Synthesis: Serine pathway was reduced from 3318 to 14 edges. Though this loses some node information, collapsing protein complexes resulted in pathways that more more closely resembled the original hypergraph in both edge distribution and topology. Collapsing protein complexes coincidentally slightly decreased class imbalance in Reactome.

After this protein-complex collapsing step, all pathways with less than 4 nodes were excluded from the analysis. This resulted in 953 Reactome pathways and 467 PathBank pathways.

Both pathway databases contain a highly skewed distribution of localizations across all interactions. The rarest localization in both databases, secretory-pathway and nucleus in Re-

ⁱ<https://github.com/indralab/pybiopax>

^j<http://www.pathwaycommons.org/pc2/formats>

actome and PathBank, respectively, occurs in less than 0.5% of all edges. The most common localization, which is cytosol for both databases, consists of 38% of Reactome interactions and 52% of PathBank interactions.

5.3.2. *Protein Localization Databases*

Two protein localization databases were used throughout all experiments, Compartments and ComPPI.

ComPPI is a meta-database for protein subcellular localizations. It combines data from 8 subcellular localization databases. It does not include data from Compartments. Proteins are assigned scores for each of 6 subcellular locations: cytosol, plasma membrane, mitochondrion, extracellular, nucleus, and secretory-pathway. These 6 locations were used for all predictions; localizations for all other data sources were mapped to these 6. ComPPI combines weights for different types of evidence across its data sources to give a probability of a protein to be found in a particular subcellular location. All human ComPPI data was retrieved on 2020-11-09.

Compartments is a protein subcellular localization database that combines data from 4 different types of data: database annotations, experimental screens, automated text mining, and predictive sequence-based models. Each data source is given a confidence score between 1 and 5 based on the level of evidence. Compartments assigns proteins to 1 of 11 subcellular locations. These 11 locations were mapped to the 6 localizations in the ComPPI database. All Compartments data was retrieved on 2021-09-29.

5.3.3. *UniProt Keyword Features*

To explore the utility of additional protein data in predicting localization data, UniProt keywords were collected for all human proteins. UniProt keywords are a controlled hierarchical vocabulary that represent a variety of protein categories such as molecular function, disease participation, structural features, and biological processes. These keywords are manually assigned and include localization data. While UniProt keywords provide a range of protein-level data, they consist of hundreds of terms, many of which are only used by a handful of proteins. Thus, they are impractical to use directly as features.

Keywords were converted into features through dimensionality reduction. Principal component analysis was performed on all keywords present in at least 5% of human proteins. Technical keywords such as “3D-Structure” and “Reference proteome” were excluded as not pertaining directly to the protein itself. Each protein was then represented by the first 6 components, chosen by a dropoff in explained variance after the first 6 components. The most important keywords for these components represented a variety of biological concepts; from functional categories such as “Tumor suppressor” and “Lipid biosynthesis” to structural features such as “ANK repeat” and “Voltage-gated channel”. These components only accounted for 42% of variance. However, given the diversity of keywords it is unlikely a small number of features could fully represent them.

5.3.4. Metrics

We used F1 score to evaluate predictive performance. The F1 score can be similarly defined as the average F1 score over all classes, treating each class as a one-against-all binary classification. The F1 score is the harmonic mean of the precision and recall. As the harmonic mean it will aggressively penalize predictions if either the precision or recall is low. It is calculated as:

$$F1 = \frac{1}{L} \left(\sum_{i=1}^L \frac{2P_i R_i}{P_i + R_i} \right)$$

where P_i is the precision of class i , calculated as:

$$P_i = \frac{TP_i}{TP_i + FP_i}$$

and R_i is the recall of class i , calculated as:

$$R_i = \frac{TP_i}{TP_i + FN_i}$$

where FP_i is the number of instances incorrectly predicted to be class i or false positives, and TP_i and FN_i are defined as above. This is considered the ‘macro’ F1 score, a ‘micro’ F1 score can also be calculated which favors the majority class.

5.4. Results

5.4.1. Pathway Localization Prediction

The Compartments database has more disagreement with pathway databases than ComPPI. While the range of ComPPI scores for all interaction localizations is wide, including a significant number of proteins existing in localizations where they have a score of 0, in all subcellular locations but secretory-pathways the median ComPPI score is highest for the corresponding Reactome localization. However, in Compartments the cytosol has the highest median score across the majority of all Reactome localizations.

Reactome and PathBank also have some disagreement in their localizations. On average, between highly matching pathways Reactome and Pathbank have 79% agreement in their assigned interaction localizations. Here, highly matching pathways were calculated as those with at least 90% of edges in one pathway appearing in the other. This shows a moderate amount of disagreement even between manually curated pathway databases.

Overall, models without any method to transfer information across a pathway, the logistic regression, random forest, and fully connected neural network models, tended to undersmooth within each pathway. The distributions for the RF, Logit, and NaivePGM models are right-skewed as compared to the true distribution, with a sizable proportion of pathways being predicted to have 5 or 6 different localizations. This is unsurprising, as these models contained no topological information with which to encourage proteins belonging to the same pathway to

have the same localization. These models greatly underestimated the proportion of pathways with a single localization.

The TrainedPGM in particular tended to oversmooth. It predicted almost all pathways as having a single localization across both datasets. It also performs particularly poorly for pathways with 2 and 3 unique localizations. This is likely due to these pathways being more evenly split between multiple localizations than those with 4 or 5 localizations, resulting in a single localization prediction to perform poorly on these pathways.

Supplementary Figures

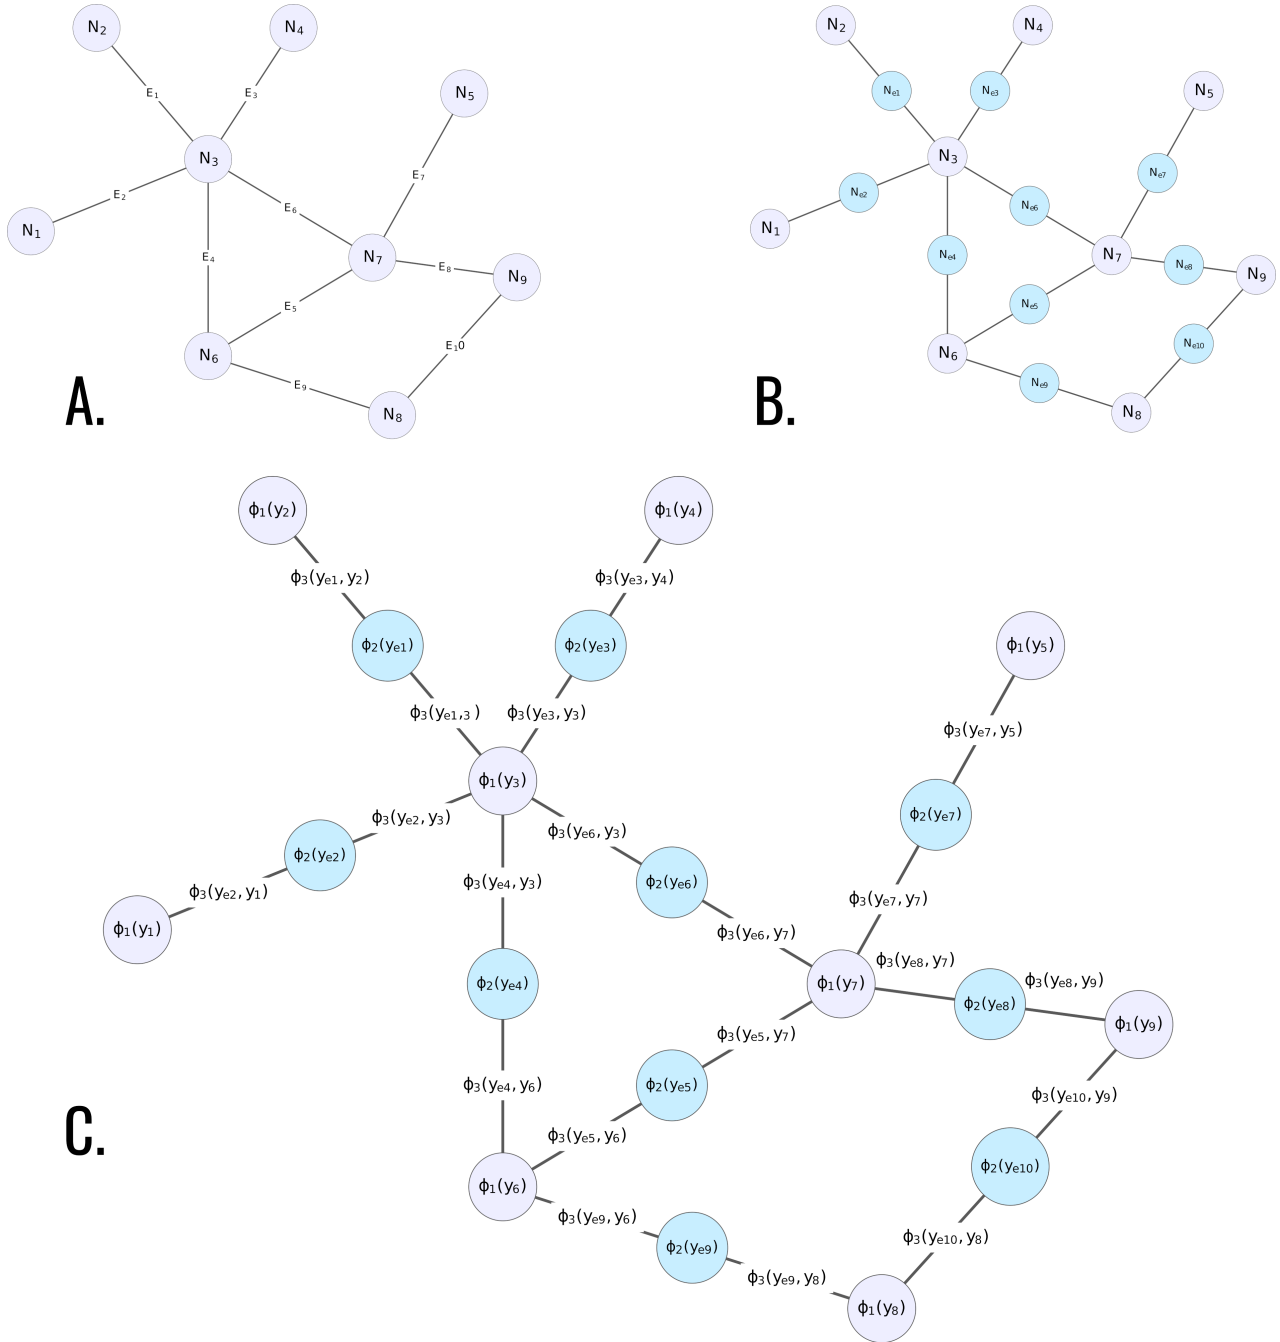


Fig. S1. Overview of how pathways were represented as probabilistic graphical models for interaction classification. Panel A shows the original pathway structure. Panel B shows then interactions-nodes that are added to each pathway. Finally, Panel C shows how potential functions are used and tied. There are 2 sets of unary potentials, $\phi_1()$ and $\phi_2()$, which model the original nodes and the interaction nodes, respectively. $\phi_3()$ models how each interaction interacts with its adjacent nodes.

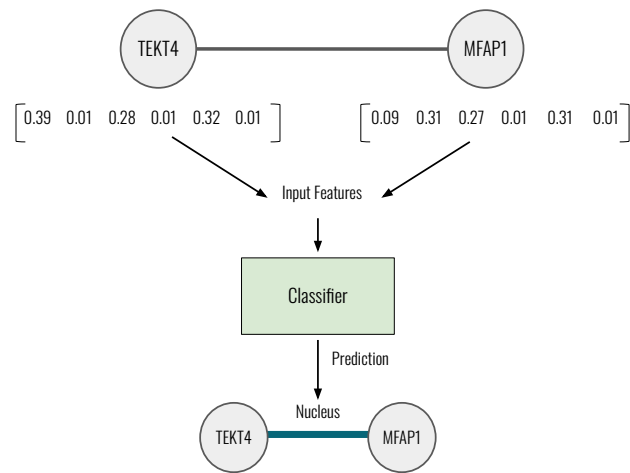


Fig. S2. Overview of how topology-free classifiers are used for the edge labeling task of localization prediction.




Fig. S3. Distribution of ComPPI protein scores by the localization of Reactome edges they belong to. Scores the probability of a protein being in a given subcellular location retrieved from the ComPPI database.




Fig. S4. Distribution of Compartments protein scores by the localization of Reactome edges they belong to. Scores are confidence scores of a protein being in a given subcellular location, weighted by the type and amount of evidence available. Scores are retrieved from the Compartments database.




Fig. S5. F1 score of predictive performance on PathBank localizations. All pathway edges are merged and measured together, resulting in 97,792 edges total.




Fig. S6. F1 score of predictive performance on Reactome localizations. All pathway edges are merged and measured together, resulting in 83,855 edges total.




Fig. S7. Distribution of the number of unique localizations in each pathway database, and as predicted by each model on each pathway database.

Supplementary Tables

Table S1: Hyperparameter ranges searched for each model.

Model	Parameter	Description/Notes	Range
All Neural Networks	Learning Rate	Learning rate for training.	$10^{-5} - 0.01$
	Linear Depth	The number of linear layers.	1 – 5
	Convolutional Depth	The number of convolutional layers. Not used by linear network.	1 – 10
	Dim	The number of dimensions in hidden layers.	24 – 128
	Dropout	Whether or not to add dropout while training.	True, False
Graph Convolutional Network		No Unique Parameters	
Graph Attention Network	Heads	The number of attention heads	1 – 5
Graph Isomorphism Network		No Unique Parameters	
Fully connected Network	Activation	Activation function used	Tanh, ReLU
Random Forest	max_depth	Maximum tree depth.	1 – 10
	min_samples_split	Minimum samples to create branches.	2 – 10
	n_estimators	Number of trees.	1 – 100
	class_weight	Whether to balance class weights.	True, False
Logistic Regression	tol	Tolerance for training.	$10^{-6} - 0.1$
	penalty	Regularization penalty to use.	L2, None
	C	Regularization strength (lower is stronger).	0.01 – 100
	class_weight	Whether to balance class weights.	True, False

Table S2: All parameter values used.

Model	Dataset	Features	Parameter	Value
FullyConnectedNN	Reactome	ComPPI		
			lRate	0.002
			l.depth	2
			dropout	0.500
			dim	83
			activation	tanh
FullyConnectedNN	Reactome	Compartments		
			lRate	0.008
			l.depth	1
			dropout	0.500
			dim	82
			activation	tanh
FullyConnectedNN	Reactome	Uniprot KW		

Continued on next page

Table S2: All parameter values used. (Continued)

Model	Dataset	Features	Parameter	Value
			lRate	$8.38e - 04$
			l.depth	3
			dropout	0.500
			dim	98
			activation	relu
FullyConnectedNN	PathBank	ComPPI		
			lRate	0.010
			l.depth	1
			dropout	0
			dim	41
			activation	tanh
FullyConnectedNN	PathBank	Compartments		
			lRate	0.009
			l.depth	2
			dropout	0
			dim	103
			activation	relu
FullyConnectedNN	PathBank	Uniprot KW		
			lRate	0.006
			l.depth	5
			dropout	0.500
			dim	81
			activation	tanh
GCN	Reactome	ComPPI		
			lRate	0.004
			l.depth	1
			dropout	0
			dim	81
			c_depth	6

Continued on next page

Table S2: All parameter values used. (Continued)

Model	Dataset	Features	Parameter	Value
GCN	Reactome	Compartments		
			lRate	$1.48e - 04$
			l_depth	1
			dropout	0.500
			dim	108
			c_depth	2
GCN	Reactome	Uniprot KW		
			lRate	0.003
			l_depth	4
			dropout	0
			dim	87
			c_depth	3
GCN	PathBank	ComPPI		
			lRate	0.002
			l_depth	1
			dropout	0
			dim	44
			c_depth	8
GCN	PathBank	Compartments		
			lRate	0.006
			l_depth	2
			dropout	0
			dim	96
			c_depth	1
GCN	PathBank	Uniprot KW		
			lRate	0.003
			l_depth	3
			dropout	0
			dim	108

Continued on next page

Table S2: All parameter values used. (Continued)

Model	Dataset	Features	Parameter	Value
			c_depth	5
GAT	Reactome	ComPPI		
			lRate	0.003
			l_depth	2
			dropout	0
			dim	48
			c_depth	7
			num_heads	4
GAT	Reactome	Compartments		
			lRate	$2.78e - 04$
			l_depth	4
			dropout	0.500
			dim	46
			c_depth	1
			num_heads	5
GAT	Reactome	Uniprot KW		
			lRate	0.010
			l_depth	2
			dropout	0
			dim	31
			c_depth	4
			num_heads	1
GAT	PathBank	ComPPI		
			lRate	0.003
			l_depth	4
			dropout	0.500
			dim	45
			c_depth	4
			num_heads	3

Continued on next page

Table S2: All parameter values used. (Continued)

Model	Dataset	Features	Parameter	Value
GAT	PathBank	Compartments		
			lRate	0.002
			l_depth	3
			dropout	0.500
			dim	43
			c_depth	3
			num_heads	4
GAT	PathBank	Uniprot KW		
			lRate	0.004
			l_depth	1
			dropout	0
			dim	39
			c_depth	2
			num_heads	4
GIN	Reactome	ComPPI		
			lRate	0.005
			l_depth	2
			dropout	0
			dim	95
			c_depth	3
GIN	Reactome	Compartments		
			lRate	$9.77e - 04$
			l_depth	4
			dropout	0.500
			dim	24
			c_depth	1
GIN	Reactome	Uniprot KW		
			lRate	0.001
			l_depth	3

Continued on next page

Table S2: All parameter values used. (Continued)

Model	Dataset	Features	Parameter	Value
			dropout	0.500
			dim	60
			c_depth	2
GIN	PathBank	ComPPI		
			lRate	0.002
			l_depth	4
			dropout	0
			dim	68
			c_depth	1
GIN	PathBank	Compartments		
			lRate	0.001
			l_depth	3
			dropout	0.500
			dim	102
			c_depth	3
GIN	PathBank	Uniprot KW		
			lRate	$9.86e - 04$
			l_depth	2
			dropout	0
			dim	80
			c_depth	1
Logit	Reactome	ComPPI		
			C	0.620
			class_weight	balanced
			penalty	l2
			tol	$1.00e - 06$
Logit	Reactome	Compartments		
			C	91.893
			class_weight	balanced

Continued on next page

Table S2: All parameter values used. (Continued)

Model	Dataset	Features	Parameter	Value
			penalty	l2
			tol	0.062
Logit	Reactome	Uniprot KW		
			C	7.604
			class_weight	balanced
			penalty	l2
			tol	0.078
Logit	PathBank	ComPPI		
			C	0.044
			class_weight	balanced
			penalty	l2
			tol	$1.84e - 05$
Logit	PathBank	Compartments		
			C	0.451
			class_weight	balanced
			penalty	l2
			tol	$4.71e - 06$
Logit	PathBank	Uniprot KW		
			C	33.176
			class_weight	balanced
			penalty	l2
			tol	0.033
RF	Reactome	ComPPI		
			class_weight	balanced
			max_depth	3
			min_samples_split	2
			n_estimators	76
RF	Reactome	Compartments		
			class_weight	balanced

Continued on next page

Table S2: All parameter values used. (Continued)

Model	Dataset	Features	Parameter	Value
			max_depth	9
			min_samples_split	3
			n_estimators	81
RF	Reactome	Uniprot KW		
			class_weight	balanced
			max_depth	6
			min_samples_split	3
			n_estimators	100
RF	PathBank	ComPPI		
			class_weight	balanced
			max_depth	10
			min_samples_split	10
			n_estimators	72
RF	PathBank	Compartments		
			class_weight	balanced
			max_depth	10
			min_samples_split	10
			n_estimators	92
RF	PathBank	Uniprot KW		
			class_weight	balanced
			max_depth	10
			min_samples_split	10
			n_estimators	58